

# Dynamic Load Balancing for Particle Methods

K. Puri, P. Ramachandran and P. Godbole

Department of Aerospace Engineering,  
IIT Bombay



Particles 2013, Stuttgart, Germany

# Outline

## 1 Particle Methods

- The serial algorithm
- The parallel algorithm
- Load Balancing

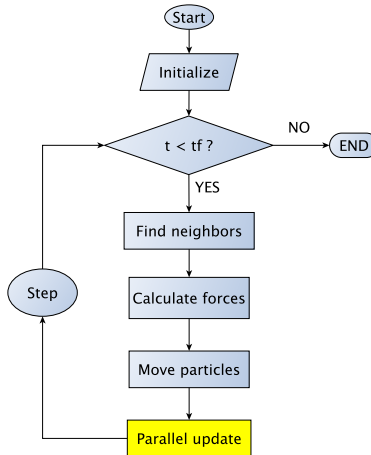
## 2 Load-balancing

- Load Balancing techniques
- Algorithms
- How do we do it?

## 3 Results

- Dam break
- Lid Driven Cavity
- Elastic collision
- Conclusion and Further work

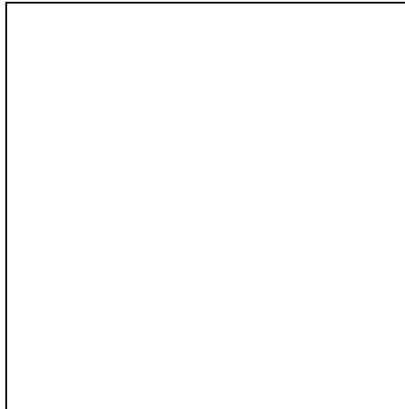
# PySPH



# Serial algorithm in a nutshell

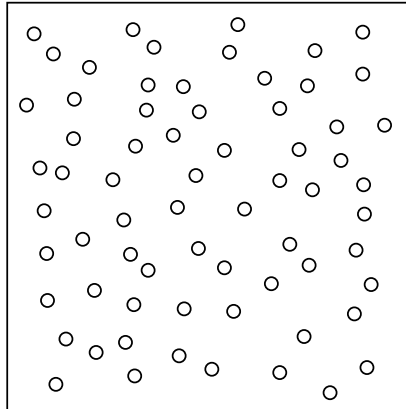
# Serial algorithm in a nutshell

Given a domain...



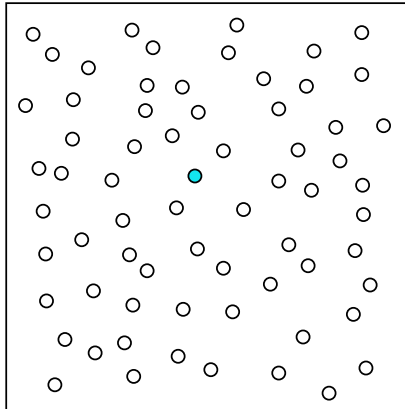
# Serial algorithm

Discretized with *Particles*



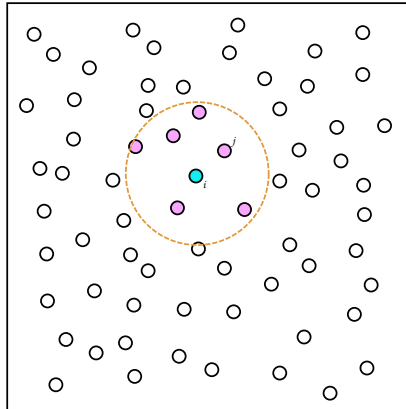
# Serial algorithm

For every particle...



# Serial algorithm

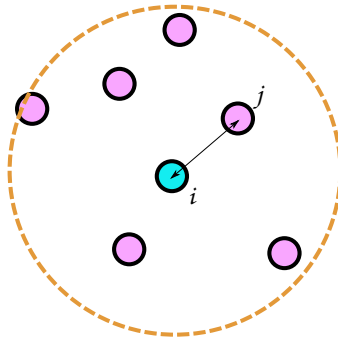
Find nearest neighbors...





# Serial algorithm

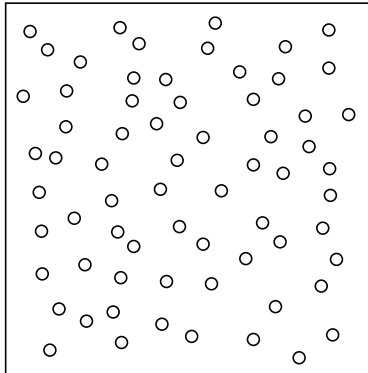
Compute interactions..



$$\frac{dU_i}{dt} = - \sum_{j \in \mathcal{N}(i)} m_j \mathcal{F}_{ij} \nabla_i W_{ij}$$

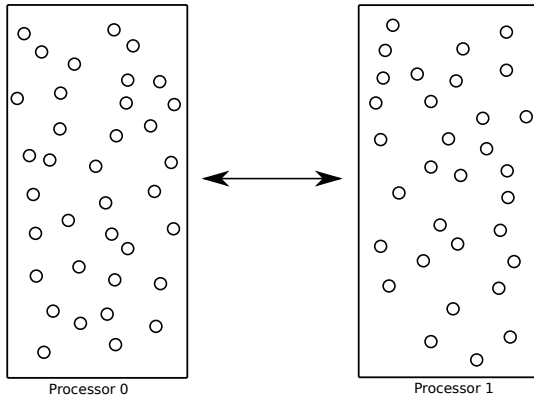
# Doing it in Parallel

Given the domain discretized with *Particles*..



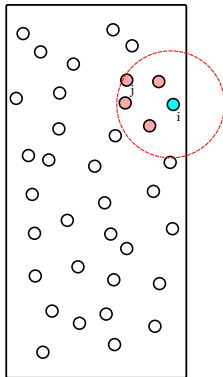
# Doing it in Parallel

Partition it across processors : *Load-balancing*

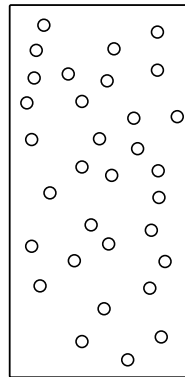


# Doing it in Parallel

For every particle, find neighbors.



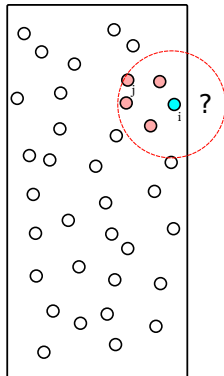
Processor 0



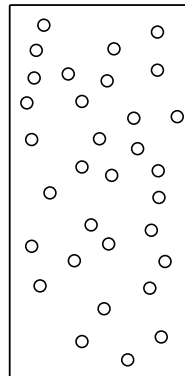
Processor 1

# Doing it in Parallel

For every particle, find neighbors. *Oops!*



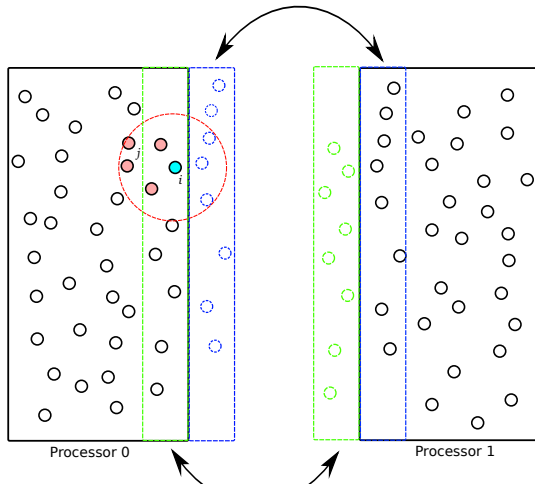
Processor 0



Processor 1

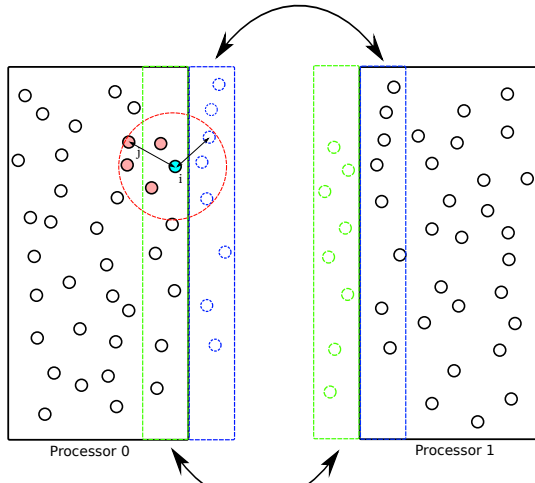
# Doing it in Parallel

Exchange ghost data.



# Doing it in Parallel

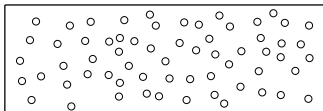
Exchange ghost data. And compute interactions..



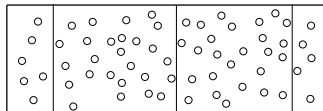
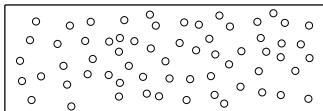
# Requirements of a good partition



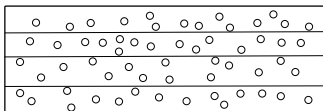
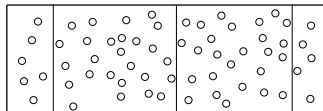
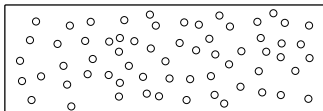
# Requirements of a good partition



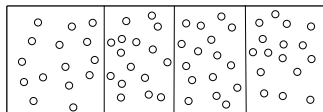
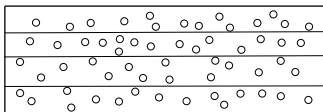
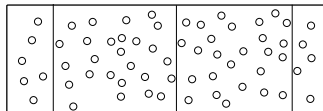
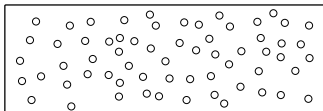
# Requirements of a good partition



# Requirements of a good partition



# Requirements of a good partition



# Outline

## 1 Particle Methods

- The serial algorithm
- The parallel algorithm
- Load Balancing

## 2 Load-balancing

- Load Balancing techniques
- Algorithms
- How do we do it?

## 3 Results

- Dam break
- Lid Driven Cavity
- Elastic collision
- Conclusion and Further work

# Scope of this work

## Algorithms

- Geometric Partitioning Algorithms
- Recursive Coordinate Bisection (RCB)
- Recursive Inertial Bisection (RIB)
- Hilbert Space Filling Curves (HSFC)

# Scope of this work

## Algorithms

- Geometric Partitioning Algorithms
- Recursive Coordinate Bisection (RCB)
- Recursive Inertial Bisection (RIB)
- Hilbert Space Filling Curves (HSFC)

## Applications

- Free surface flows
- Fixed domain incompressible NS
- Elastic solid mechanics

# Scope of this work

## Algorithms

- Geometric Partitioning Algorithms
- Recursive Coordinate Bisection (RCB)
- Recursive Inertial Bisection (RIB)
- Hilbert Space Filling Curves (HSFC)

## Applications

- Free surface flows
- Fixed domain incompressible NS
- Elastic solid mechanics

- Partitioning quality
- Execution times
- Scale-up



# Algorithm classification

## Geometric partitioners

- Physical coordinates as input
- Most general for numerical work
- Natural for particle methods

# Algorithm classification

## Geometric partitioners

- Physical coordinates as input
- Most general for numerical work
- Natural for particle methods

## Examples

- Recursive Coordinate Bisection (RCB)
- Recursive Inertial Bisection (RIB)
- Space Filling Curves (SFC)

# Algorithm classification

## Graph partitioners

- Data represented as a graph
- Inherently suitable for mesh-based methods
- Cell based graph-partitioning may be used for SPH

# Algorithm classification

## Graph partitioners

- Data represented as a graph
- Inherently suitable for mesh-based methods
- Cell based graph-partitioning may be used for SPH

## Examples

- METIS/ParMETIS
- PTScotch
- Hypergraph partitioning

# Recursive Coordinate Bisection (RCB)

## Algorithm and Advantages

- Recursive
- Orthogonal cuts
- Fast

# Recursive Coordinate Bisection (RCB)

## Algorithm and Advantages

- Recursive
- Orthogonal cuts
- Fast

## Disadvantages

- Not rotationally invariant
- Stretched halo regions

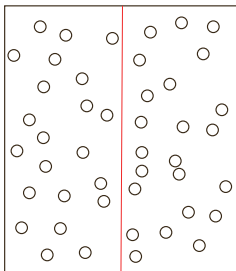
# Recursive Coordinate Bisection (RCB)

## Algorithm and Advantages

- Recursive
- Orthogonal cuts
- Fast

## Disadvantages

- Not rotationally invariant
- Stretched halo regions



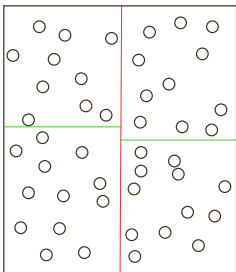
# Recursive Coordinate Bisection (RCB)

## Algorithm and Advantages

- Recursive
- Orthogonal cuts
- Fast

## Disadvantages

- Not rotationally invariant
- Stretched halo regions





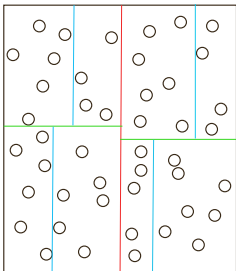
# Recursive Coordinate Bisection (RCB)

## Algorithm and Advantages

- Recursive
- Orthogonal cuts
- Fast

## Disadvantages

- Not rotationally invariant
- Stretched halo regions



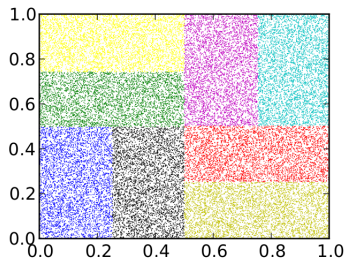
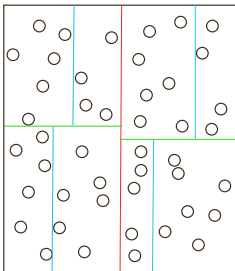
# Recursive Coordinate Bisection (RCB)

## Algorithm and Advantages

- Recursive
- Orthogonal cuts
- Fast

## Disadvantages

- Not rotationally invariant
- Leads to stretched halo-regions



# Recursive Inertial Bisection (RIB)

## Algorithm and Advantages

- Variant of RCB
- Orthogonal cuts to principal inertial axes
- Adaptive to rotations
- Lesser communication overhead

# Recursive Inertial Bisection (RIB)

## Algorithm and Advantages

- Variant of RCB
- Orthogonal cuts to principal inertial axes
- Adaptive to rotations
- Lesser communication overhead

## Disadvantages

- Eigenvector computations

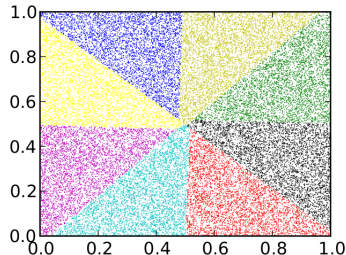
# Recursive Inertial Bisection (RIB)

## Algorithm and Advantages

- Variant of RCB
- Orthogonal cuts to principal inertial axes
- Adaptive to rotations
- Lesser communication overhead

## Disadvantages

- Eigenvector computations



# Hilbert Space Filling Curves (HSFC)

## Algorithm and Advantages

- Use a SFC  $f : R^3 \rightarrow R$
- Order objects linearly
- Geometric locality

# Hilbert Space Filling Curves (HSFC)

## Algorithm and Advantages

- Use a SFC  $f : R^3 \rightarrow R$
- Order objects linearly
- Geometric locality

## Disadvantages

- Particle distribution has *projections*
- Disconnected regions for complex geometries

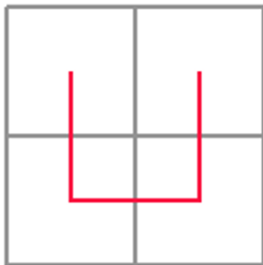
# Hilbert Space Filling Curves (HSFC)

## Algorithm and Advantages

- Use a SFC  $f : R^3 \rightarrow R$
- Order objects linearly
- Geometric locality

## Disadvantages

- Particle distribution has *projections*
- Disconnected regions for complex geometries





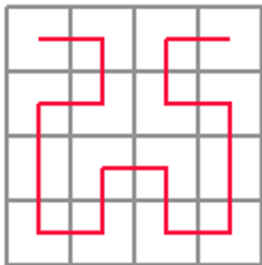
# Hilbert Space Filling Curves (HSFC)

## Algorithm and Advantages

- Use a SFC  $f : R^3 \rightarrow R$
- Order objects linearly
- Geometric locality

## Disadvantages

- Particle distribution has *projections*
- Disconnected regions for complex geometries



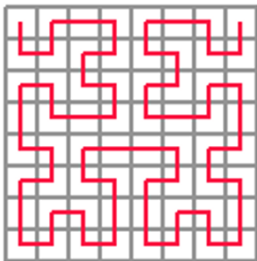
# Hilbert Space Filling Curves (HSFC)

## Algorithm and Advantages

- Use a SFC  $f : R^3 \rightarrow R$
- Order objects linearly
- Geometric locality

## Disadvantages

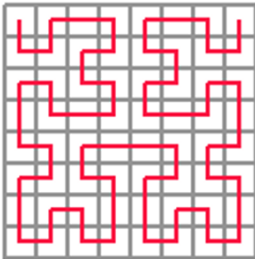
- Particle distribution has *projections*
- Disconnected regions for complex geometries



# Hilbert Space Filling Curves (HSFC)

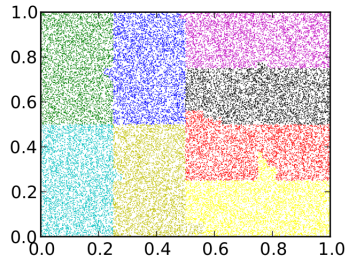
## Algorithm and Advantages

- Use a SFC  $f : R^3 \rightarrow R$
- Order objects linearly
- Geometric locality

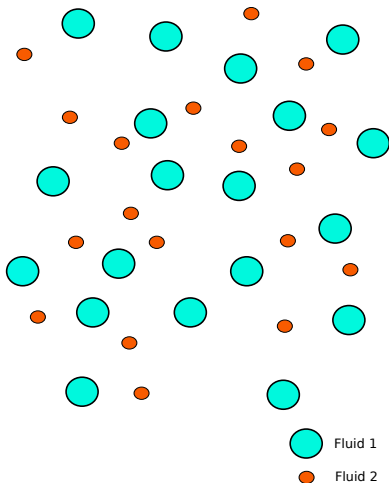


## Disadvantages

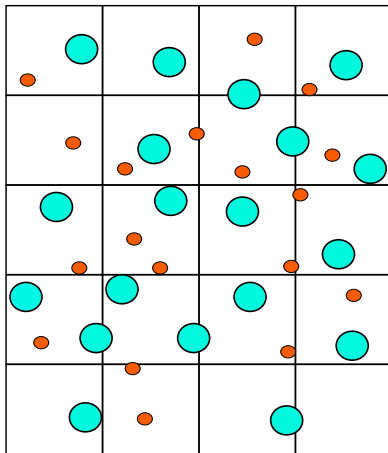
- Particle distribution has *projections*
- Disconnected regions for complex geometries



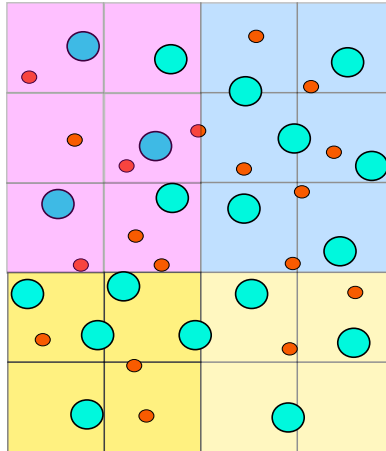
# How do we do it?



# How do we do it?



# How do we do it?



# Zoltan Data Management Library

## What is it?

- Developed by Sandia National Laboratories
- Trilinos Project (9.0 September 2008)
- Zoltan v3.6 released in September 2011

## What can it do?

- **Dynamic Load Balancing**
- Graph Coloring
- Dynamic memory management

# Outline

## 1 Particle Methods

- The serial algorithm
- The parallel algorithm
- Load Balancing

## 2 Load-balancing

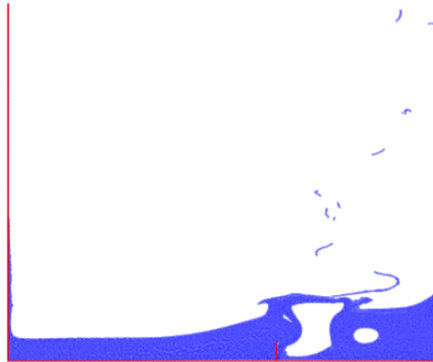
- Load Balancing techniques
- Algorithms
- How do we do it?

## 3 Results

- Dam break
- Lid Driven Cavity
- Elastic collision
- Conclusion and Further work



## 2D Dam Break

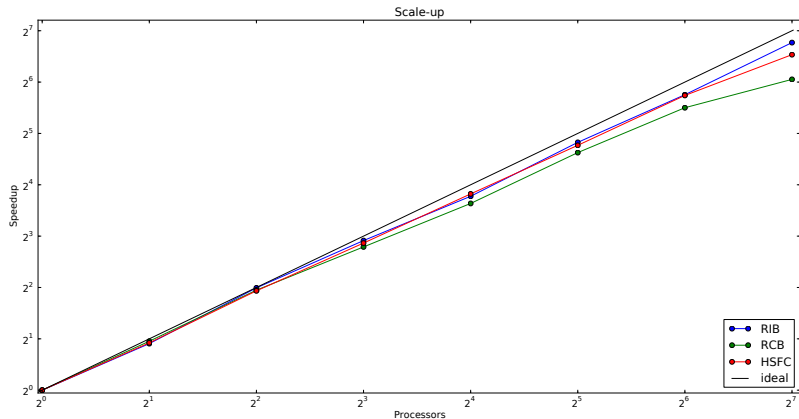


- $N_p \approx O(0.1M)$

# Machine Architecture

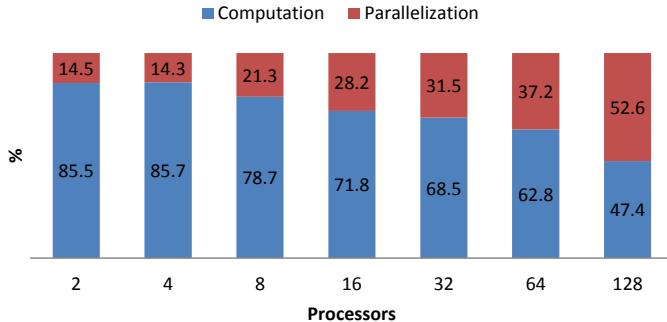
- Linux (CentOS) cluster
- Six-core AMD Opteron
- 1 Gigabit Ethernet interconnect
- 12 GB RAM per node

# Scale-up : $N_p \approx 10M$ , per-iteration



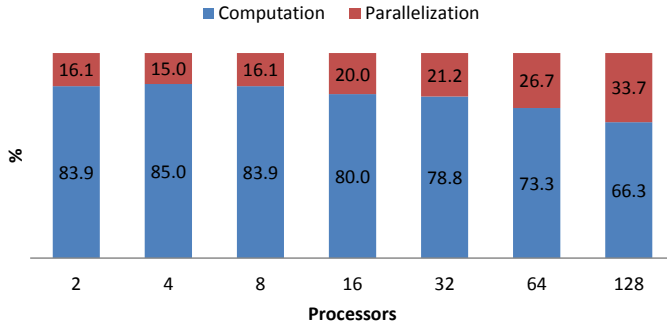
# Time distribution : RCB

## RCB Time distribution



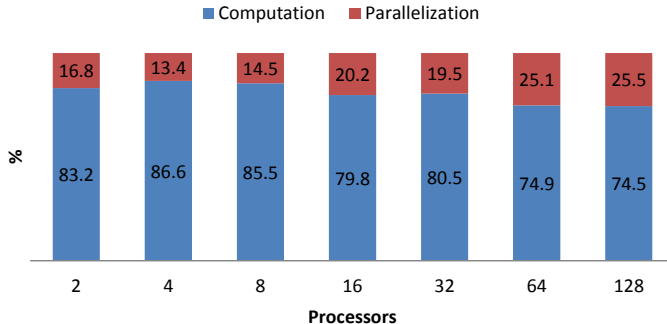
# Time distribution : HSFC

## HSFC Time distribution

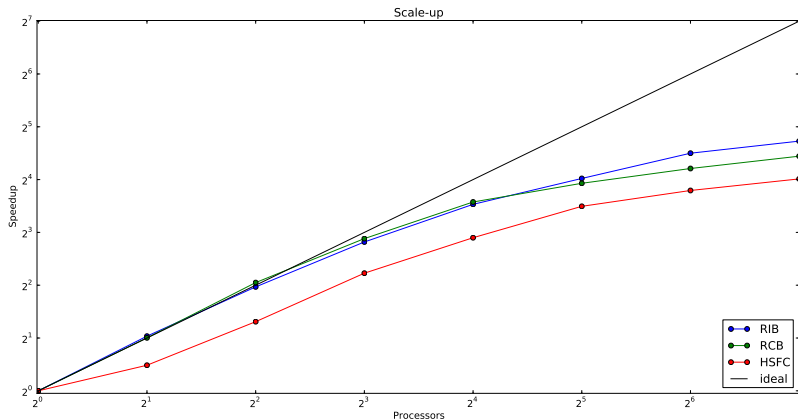


# Time distribution : RIB

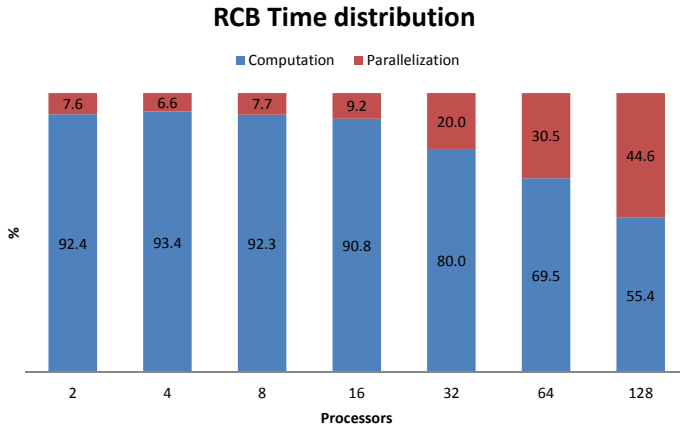
## RIB Time distribution



# 3D Dam break : Scale-up : $N_p \approx 10M$ , per-iteration

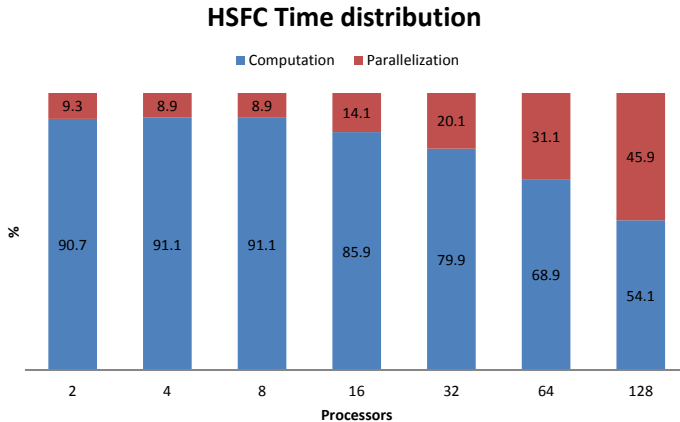


## 3D Dam break : RCB Time distribution



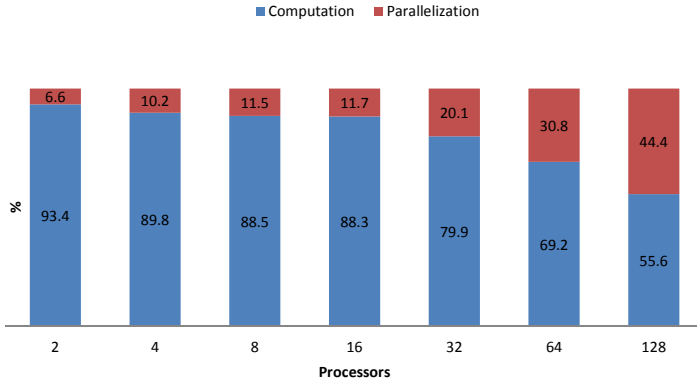


## 3D Dam break : HSFC Time distribution

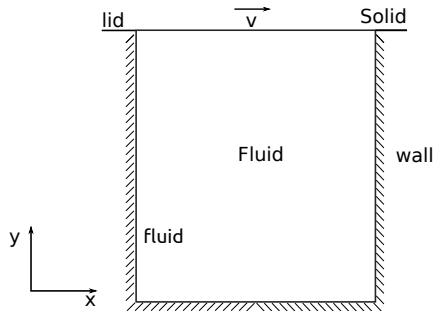


## 3D Dam break : RIB Time distribution

### RIB Time distribution

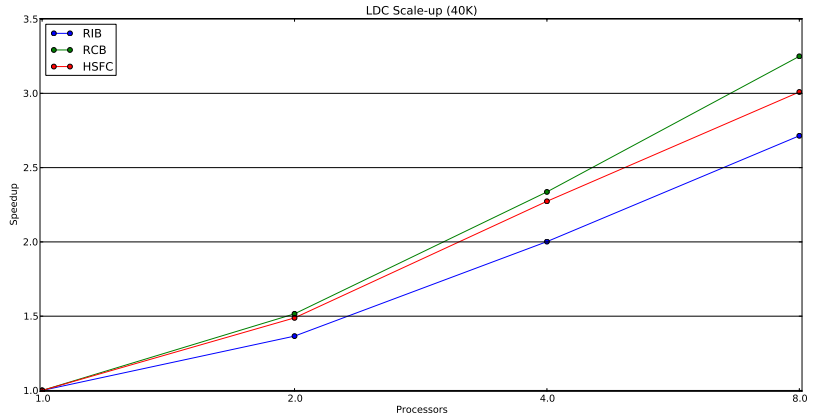


# Lid Driven Cavity

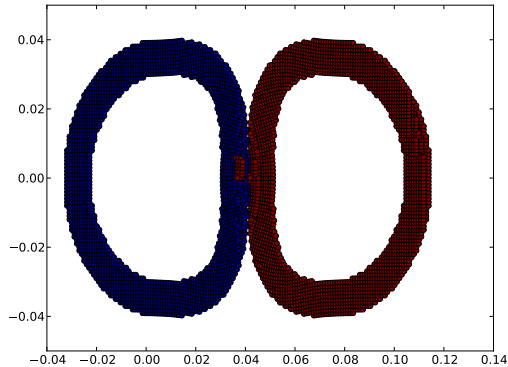


- Fixed 2D domain
- $N_p \leq 10^5$

# Scale up

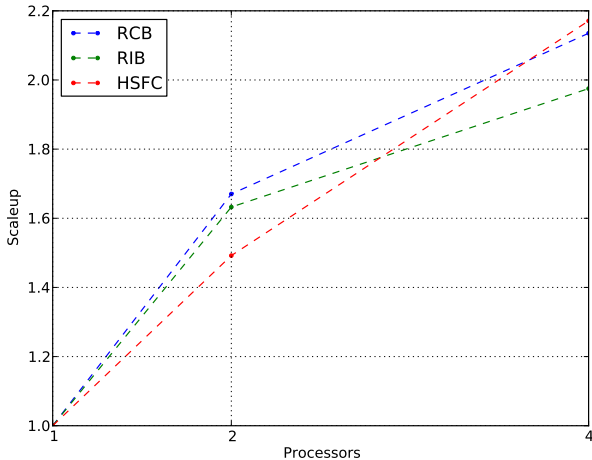


# Elastic collision



- Collision
- $N_p \approx 20,000$

# Scale up



# Conclusions

- Evaluation of geometric load balancing algorithms
- RIB for dynamic free surfaces
- RCB/HSFC for contact problems
- RCB for fixed domain problems

## Future work

- HVI and 3D contact problems
- Scale up for massively parallel computers
- Coupled methods



Thank you!